

Operating Systems

(Sistemas Operativos)

The Process Abstraction

University of Minho
2024-2025



What will we learn?

CPU Management

- ◎ Although a computer has a limited number of CPUs, users get the impression that several programs (a lot more than the available CPUs) are running simultaneously
 - How does the OS provide this illusion?
 - How can users use the OS APIs to start and stop programs?
 - How does the OS choose what programs will be running, be switched, ...?
- ◎ Let us start with the **process** abstraction!

Process

Definition

- ◎ Process: a running program¹
- ◎ The process abstraction is fundamental for virtualizing the CPU
 - In a given CPU, the OS can start a process, stop it, run another one
- ◎ With the right mechanisms (**context switch**) and policies (**scheduling algorithms**)
 - Time sharing and multiprogramming become possible!
 - But, first, one must understand the Process abstraction
- ◎ Definitions to take in account:
 - Program: passive entity (e.g., an executable file sitting on disk with several instructions)
 - Process: active entity, executing instructions and accessing system resources

¹ Informal definition, as we will see, one program can create several processes

Process


The abstraction

- Being an active entity, the process will access/use system resources
 - CPU - to execute instructions of a program
 - Memory - to keep the instructions, variables, data structures ... of a program
 - Disk - to store the program's information persistently
- The current activity of a process can be represented by
 - The value of the **Program Counter** (PC) and other **CPU registers**
 - The PC is a special register identifying the instruction to execute next
 - Many program instructions read/update CPU registers
 - **Address space**, i.e., the memory layout of a process
 - Where the instructions ("code"), data (variables, structures, ...) of a process reside
 - List of **open files**, as processes may do I/O

Process Creation

How are programs transformed into processes

- Programs (*i.e.*, code¹ and static data²) initially reside on disk as executable files
 - The source code files of programs are compiled into object files (with the aid of a **compiler**)
 - Object files, and required libraries (*e.g.*, *standard C library*), are combined into one executable binary file (with the **linker**)




¹ In the literature, the Code segment is also known as Text

² Static data (also known as the Data segment in the literature) has, for instance, initialized (global) variables

Process Creation

How are programs transformed into processes

- Programs (*i.e.*, code¹ and static data²) initially reside on disk as executable files
 - The source code files of programs are compiled into object files (with the aid of a **compiler**)
 - Object files, and required libraries (*e.g.*, *standard C library*), are combined into one executable binary file (with the **linker**)
- Bytes from the binary file (stored on disk) are loaded into memory (with the **loader**)
 - Loading can be done **eagerly** (all bytes are loaded before running the program) or **lazily** (pieces of the program are loaded as these are needed and while the program is executing)³



¹ In the literature, the Code segment is also known as Text


² Static data (also known as the Data segment in the literature) has, for instance, initialized (global) variables

³ Modern OSs use lazy loading

Process Creation

How are programs transformed into processes


- Once code and static data are loaded into memory, the OS
 - Allocates memory for the **stack** (e.g., arguments, local variables, function parameters, return addresses)
 - Allocates memory for the **heap** (dynamically allocated memory)
 - Does other tasks related to I/O (e.g., initialize default file descriptors for *stdin*, *stdout* and *stderr*)



Process Creation

How are programs transformed into processes

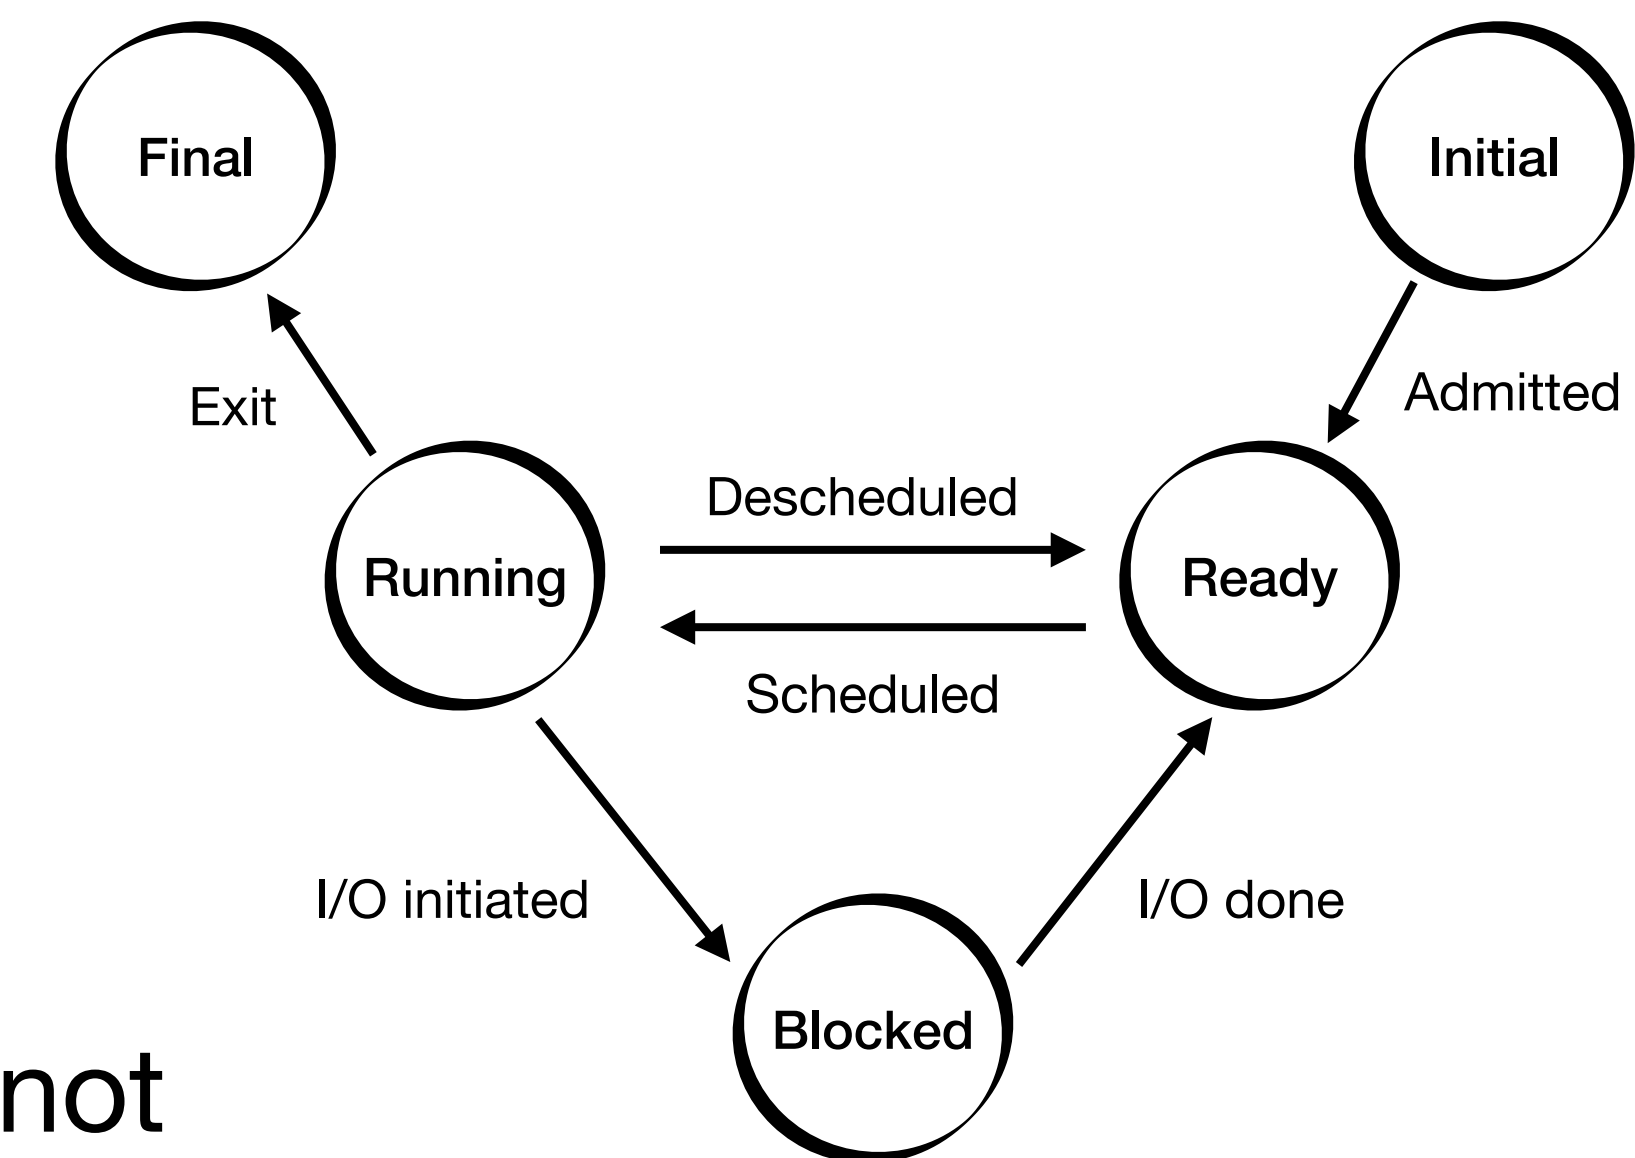
- Once code and static data are loaded into memory, the OS
 - Allocates memory for the **stack** (e.g., arguments, local variables, function parameters, return addresses)
 - Allocates memory for the **heap** (dynamically allocated memory)
 - Does other tasks related to I/O (e.g., initialize default file descriptors for *stdin*, *stdout* and *stderr*)
- Run the program at its entry point, namely the *main()* routine
 - The OS transfers control of the CPU to the newly created process to begin executing



Process Lifecycle

States

- In a simplified view a process can be
 - **Running** on the processor (executing instructions)
 - **Ready** to run (waiting for processor time)
 - **Blocked** until some other event takes place (e.g., when a process makes a disk I/O request it becomes blocked until the request is served)
- We can extend the above states with
 - **Initial** - the process is being created
 - **Final** - the process has finished execution but it was not cleaned up yet (**zombie** state in Unix)
 - Some OSs may further delineate their states



Process Lifecycle

States

● During execution, a process may transition several times through the **Running**, **Ready** and **Blocked** states

● At the example on the right

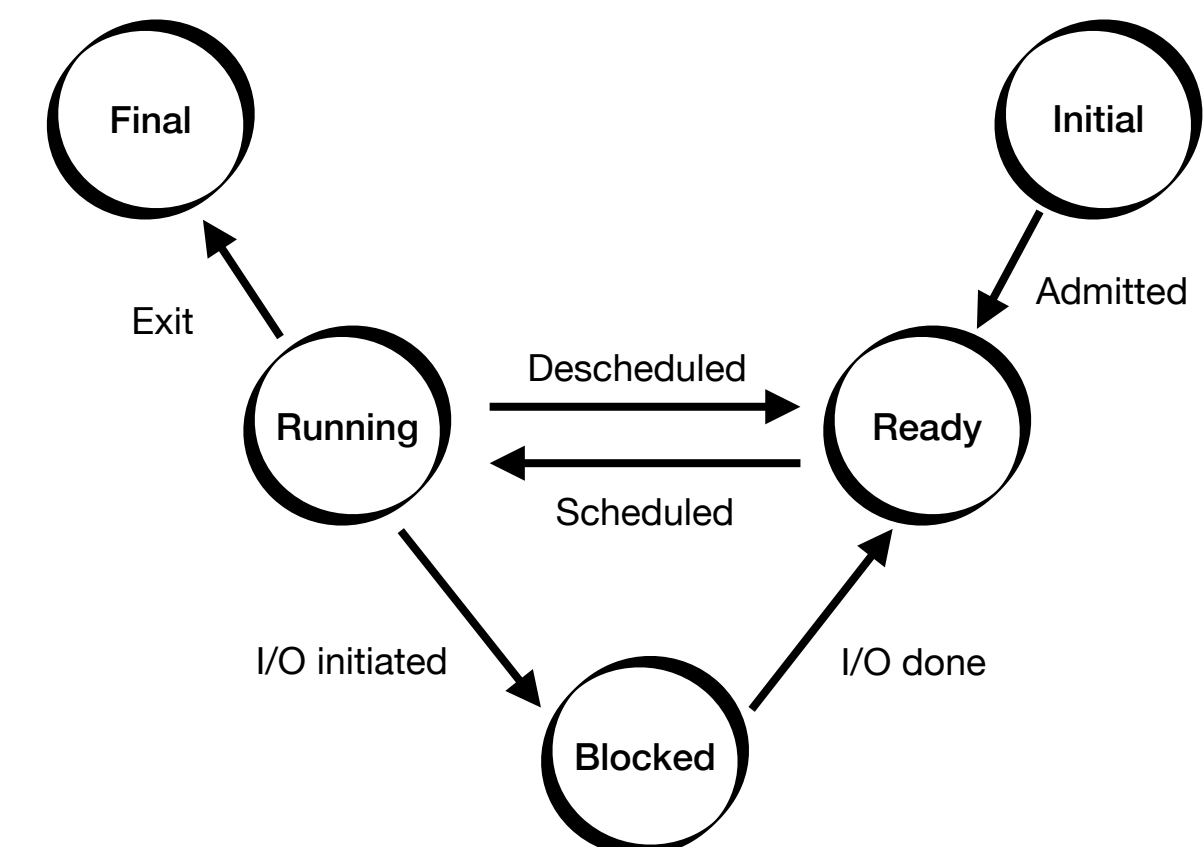
- ▶ Process₀ blocks for I/O and Process₁ is scheduled (time 4)
- ▶ Process₀ is ready but the OS decides to keep Process₁ running (time 7)

● Questions

- ▶ Why is it a good idea to schedule another process when Process₁ blocks for I/O (time 4)?
- ▶ Should one schedule immediately Process₀ when I/O is done (time 7)?

● These decisions are made by the OS **scheduler**

Time	Process ₀	Process ₁	Notes
1	Running	Ready	
2	Running	Ready	
3	Running	Ready	
4	Blocked	Running	Process ₀ initiates I/O Process ₀ is blocked, so Process ₁ runs
5	Blocked	Running	
6	Blocked	Running	
7	Ready	Running	I/O is done
8	Ready	Running	Process ₁ finishes
9	Running	Finished	
10	Running	Finished	Process ₀ finishes



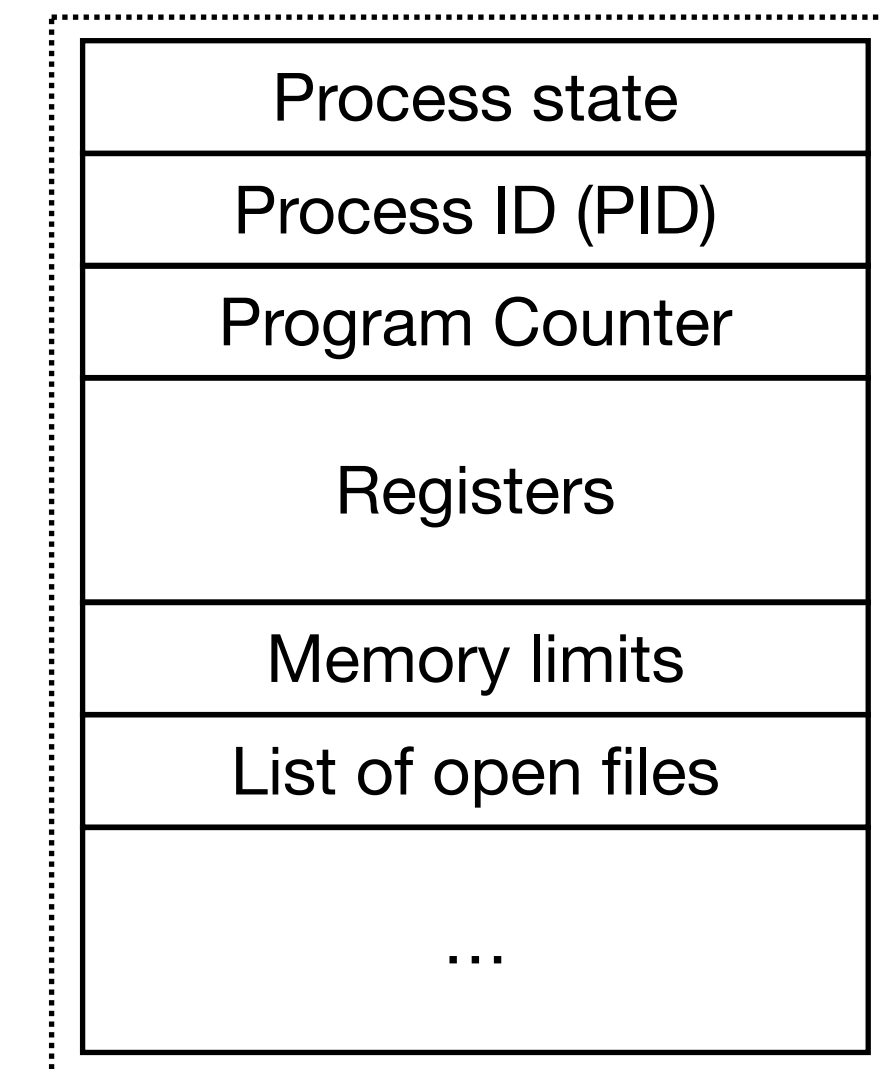
Operating System

Data Structures

- The OS is a program so it uses data structures!
 - A **process list** is used to track the state of processes (*i.e.*, **Running, Ready, Blocked**)
 - Each entry of this list is called a **Process Control Block (PCB)**

- PCB

- **Process State:** Running, Ready, Blocked, ...
- **Process ID (PID):** Unique identifier for the process
- **Program Counter:** memory address of the next instruction to execute
- **Registers:** CPU registers that must be saved in memory so the process may be resumed later (e.g., stack pointers)
- **Memory limits:** The memory limits (address space) the process may access¹
- **List of open files:** Files opened by the process¹



PCB

¹ To be further discussed at the memory management and persistence classes!

Process API

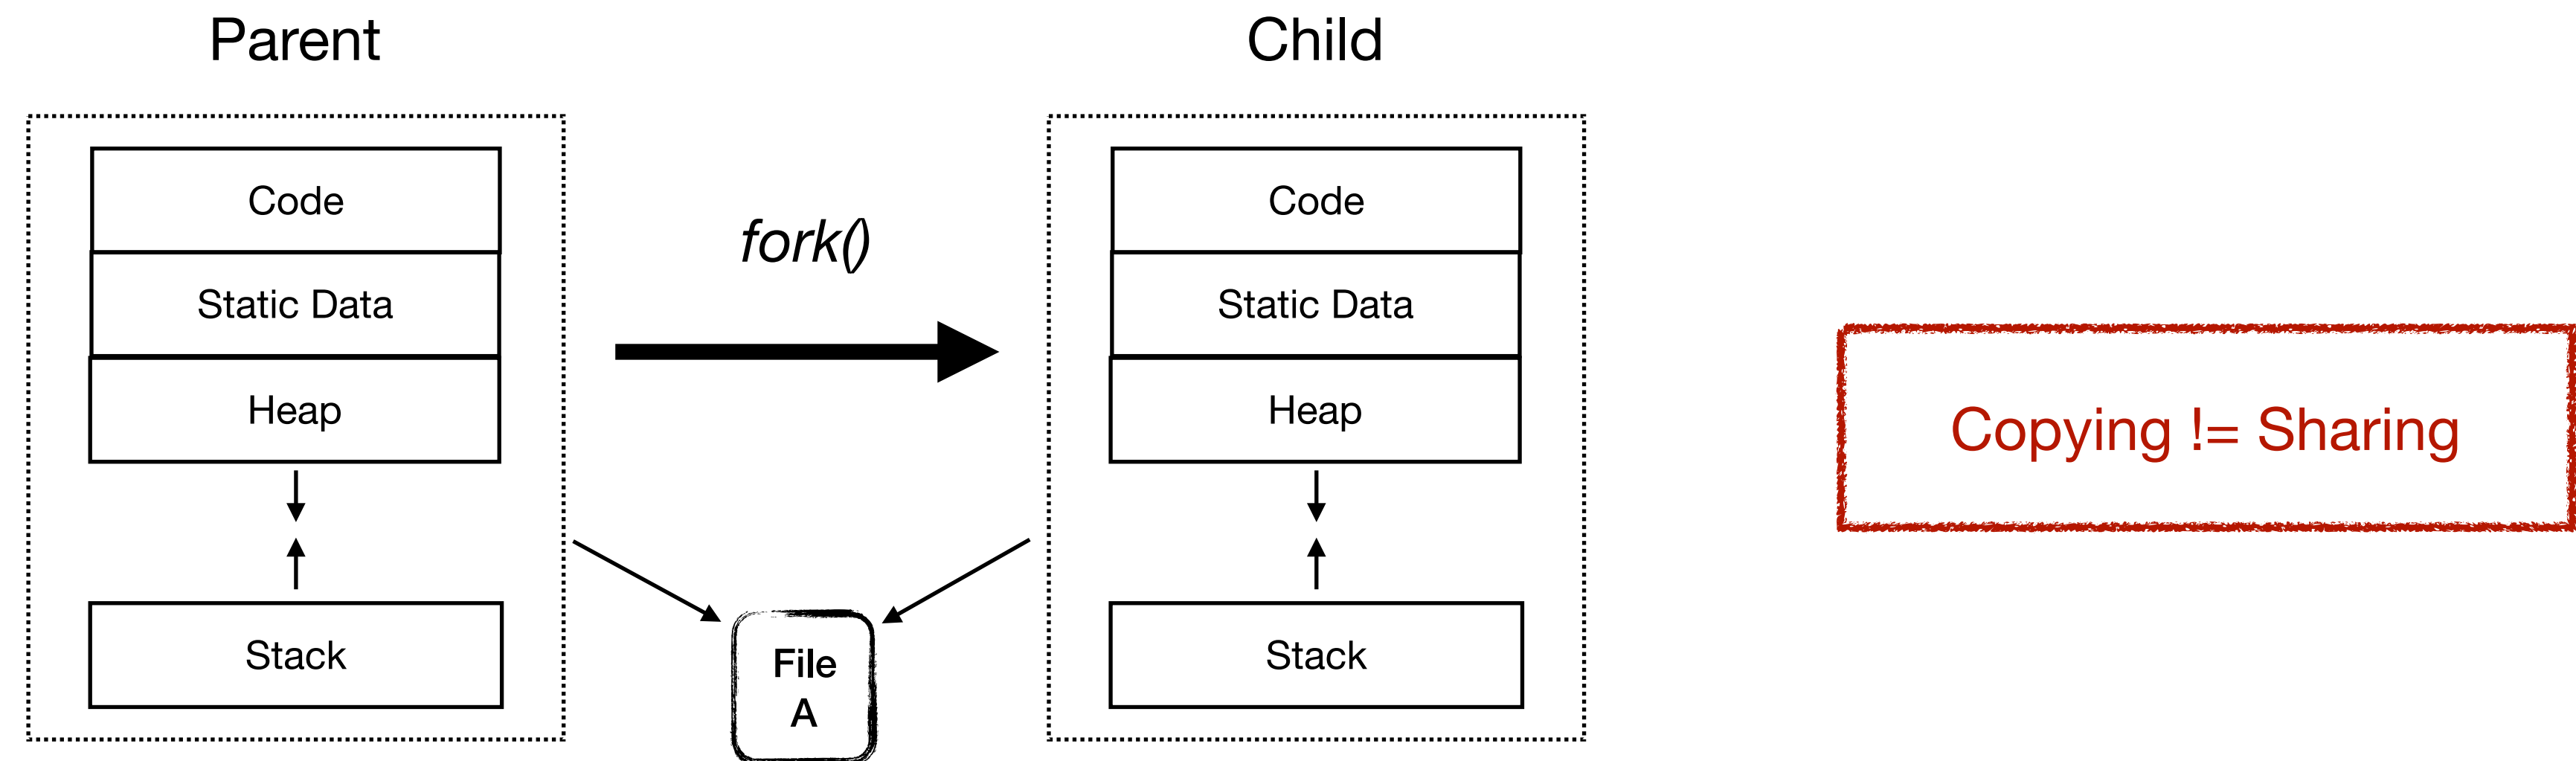
◎ The API must allow to

- **Create** new processes, otherwise, how can one launch a new program?
- **Destroy** processes, when one wishes to terminate processes non-gracefully
- **Wait** for a process until it finishes its execution
- Check the **status** of a process (e.g., its current state, execution time)
- Other controls such as **suspending** and **resuming** a process on demand

UNIX Process API

Fork

- The *fork()* system call is used to create a new process - *pid_t fork(void)*
- When fork is called by the **parent** process, the **child** process, an almost identical copy of the calling process is created
 - The **child process has its own copy** of the address space, registers, PC, list of open files




UNIX Process API

Fork

- The **value returned** by *fork()* to each process is different
 - The parent receives the **PID of the newly-created child**
 - The child receives a return code of **zero**

- Two processes are now running **concurrently**
 - Accessing an independent address space (memory)
 - **Question:** What if the two processes write to *file A*?

- Two processes are now executing the instructions of the same program
 - **Question:** How can one specify code to be executed by each process?



Fork

Example (Try it!)

Prompt

`prompt> ./p1`



```
1 #include <unistd.h>
2 #include <stdio.h>
3
4 int main(int argc, char* argv[]){
5
6     printf("Hello (pid:%d)\n", getpid());
7     pid_t rc_fork = fork();
8     if(rc_fork == 0){
9         printf("Child (pid: %d)\n", getpid());
10        // Do stuff (child)
11    }else{
12        printf("Parent of %d (pid: %d)\n", rc_fork, getpid());
13        //Do stuff (parent)
14    }
15
16    return 0;
17 }
```

© Try compiling and running this code!

Fork

Example (Try it!)

Process 44562

Prompt



```
1 #include <unistd.h>
2 #include <stdio.h>
3
4 int main(int argc, char* argv[]){
5
6 → printf("Hello (pid:%d)\n", getpid());
7     pid_t rc_fork = fork();
8     if(rc_fork == 0){
9         printf("Child (pid: %d)\n", getpid());
10        // Do stuff (child)
11    }else{
12        printf("Parent of %d (pid: %d)\n", rc_fork, getpid());
13        //Do stuff (parent)
14    }
15
16    return 0;
17 }
```

```
prompt> ./p1
Hello (pid:44562)
```

© When the program starts, the process prints the message “Hello” and its PID - 44562

Fork

Example (Try it!)

Process 44562
(parent)

Process 44563
(child)

Prompt



```
1 #include <unistd.h>
2 #include <stdio.h>
3
4 int main(int argc, char* argv[]){
5
6     printf("Hello (pid:%d)\n", getpid());
7     pid_t rc_fork = fork();
8     if(rc_fork == 0){
9         printf("Child (pid: %d)\n", getpid());
10        // Do stuff (child)
11    }else{
12        printf("Parent of %d (pid: %d)\n", rc_fork, getpid());
13        //Do stuff (parent)
14    }
15
16    return 0;
17 }
```



```
1 #include <unistd.h>
2 #include <stdio.h>
3
4 int main(int argc, char* argv[]){
5
6     printf("Hello (pid:%d)\n", getpid());
7     pid_t rc_fork = fork();
8     if(rc_fork == 0){
9         printf("Child (pid: %d)\n", getpid());
10        // Do stuff (child)
11    }else{
12        printf("Parent of %d (pid: %d)\n", rc_fork, getpid());
13        //Do stuff (parent)
14    }
15
16    return 0;
17 }
```

```
prompt> ./p1
Hello (pid:44562)
```

- When *fork()* is called a new process (child) is created, which is an almost identical copy of the calling process (parent)

Fork

Example (Try it!)

Process 44562
(parent)

Process 44563
(child)

Prompt

```
1 #include <unistd.h>
2 #include <stdio.h>
3
4 int main(int argc, char* argv[]){
5
6     printf("Hello (pid:%d)\n", getpid());
7     pid_t rc_fork = fork();
8     if(rc_fork == 0){
9         printf("Child (pid: %d)\n", getpid());
10        // Do stuff (child)
11    }else{
12        → printf("Parent of %d (pid: %d)\n", rc_fork, getpid());
13        //Do stuff (parent)
14    }
15
16    return 0;
17 }
```

```
1 #include <unistd.h>
2 #include <stdio.h>
3
4 int main(int argc, char* argv[]){
5
6     printf("Hello (pid:%d)\n", getpid());
7     pid_t rc_fork = fork();
8     if(rc_fork == 0){
9         → printf("Child (pid: %d)\n", getpid());
10        // Do stuff (child)
11    }else{
12        printf("Parent of %d (pid: %d)\n", rc_fork, getpid());
13        //Do stuff (parent)
14    }
15
16    return 0;
17 }
```

```
prompt> ./p1
Hello (pid:44562)
Parent of 44563 (pid: 44562)
Child (pid: 44563)
prompt>
```

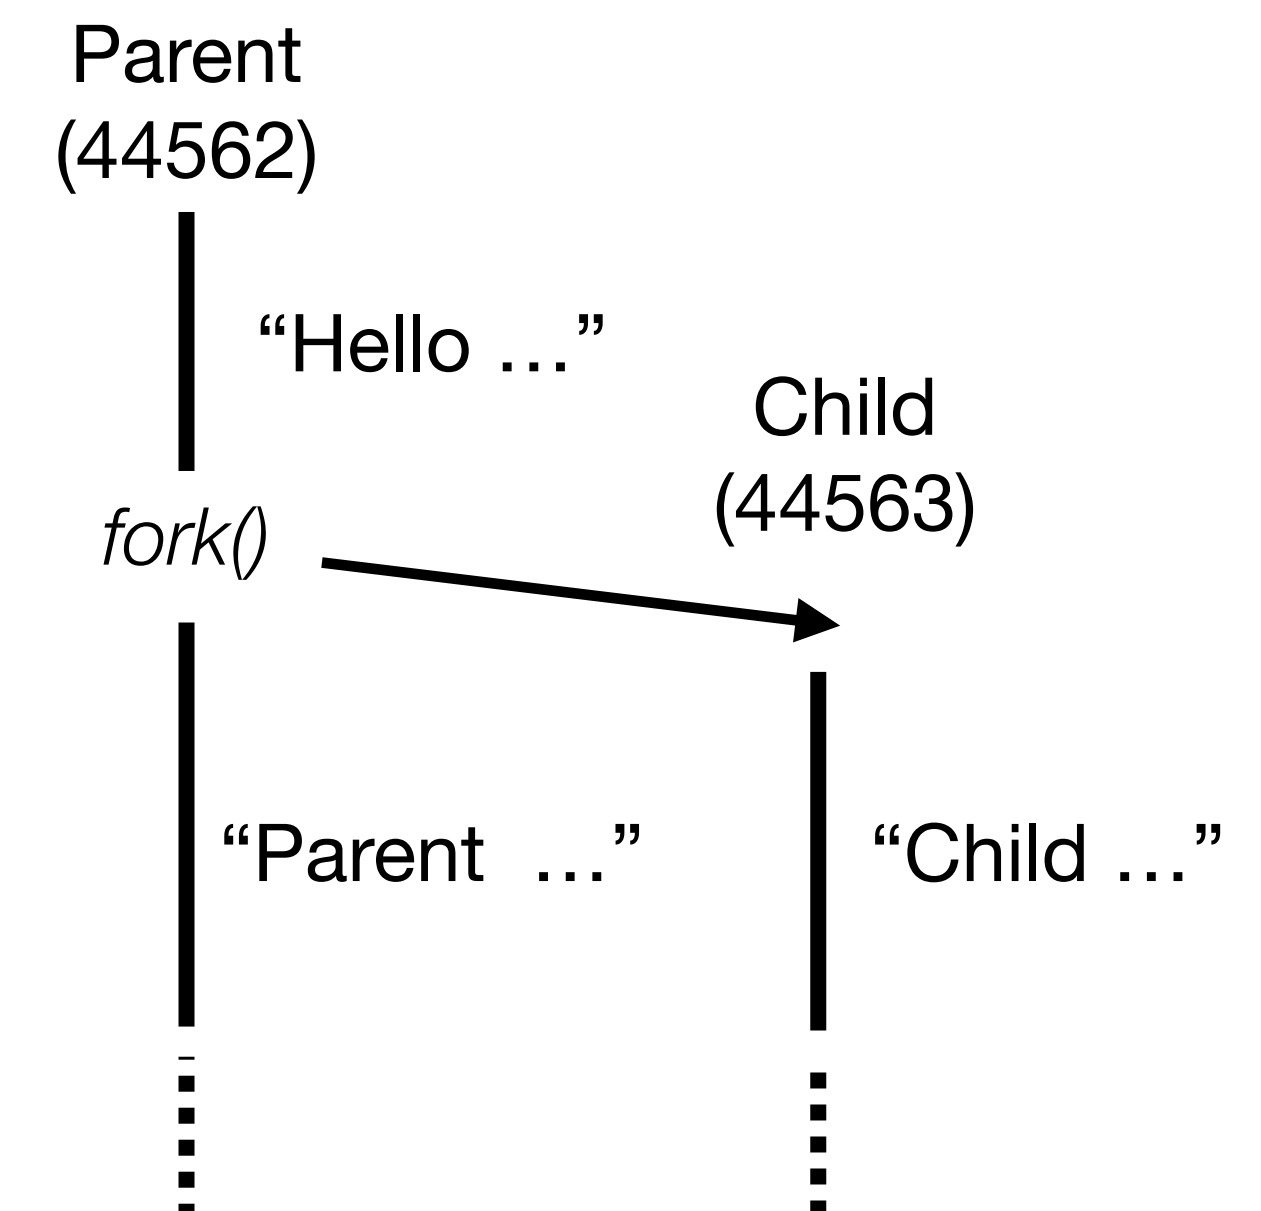
- The child does not start running at *main()*, it starts as if it had called *fork()* itself (“Hello” is only printed once)
- The child is not an identical copy since it receives a different return code from *fork()*.
Useful for specifying different code (doing different actions) for each process

Fork

Concurrency

- **Question:** If the previous program runs multiple times, the messages “*Parent ...*” and “*Child ...*” may be switched. Why?
- **Question:** even assuming a single CPU, the output is non deterministic, why?
- Non-determinism leads to interesting **concurrency** problems
 - Although the memory of each process is isolated (private), processes may share access to other resources (e.g., files)
 - Memory is not isolated for multi-threaded programs
 - **Threads** of a process have their own program counter but share the process’s address space

```
prompt> ./p1
Hello (pid:44562)
Parent of 44563 (pid: 44562)
Child (pid: 44563)
prompt>
```



Fork

Another example

◎ **Question:** And with this code...

- ▶ How many times is the message “*I’m Here*” printed?
- ▶ What process(es) print that message?
- ▶ What messages can be switched?

```
1 int main(int argc, char* argv[]){
2
3     printf("Hello (pid:%d)\n", getpid());
4     pid_t rc_fork = fork();
5     if(rc_fork == 0){
6         printf("Child (pid: %d)\n", getpid());
7         // Do stuff (child)
8     }else{
9         printf("Parent of %d (pid: %d)\n", rc_fork, getpid());
10        //Do stuff (parent)
11    }
12    }
13    printf("I'm here (pid: %d)\n", getpid());
14    return 0;
15 }
```

UNIX Process API

Wait

- What if one wants the parent to wait for its child to finish executing?
- The *wait()* system call blocks the execution of the calling process until one of the child processes ends - *pid_t wait(int *status);*
 - **Disclaimer:** Some variants of *wait()* may not block until the child process exits

Wait example

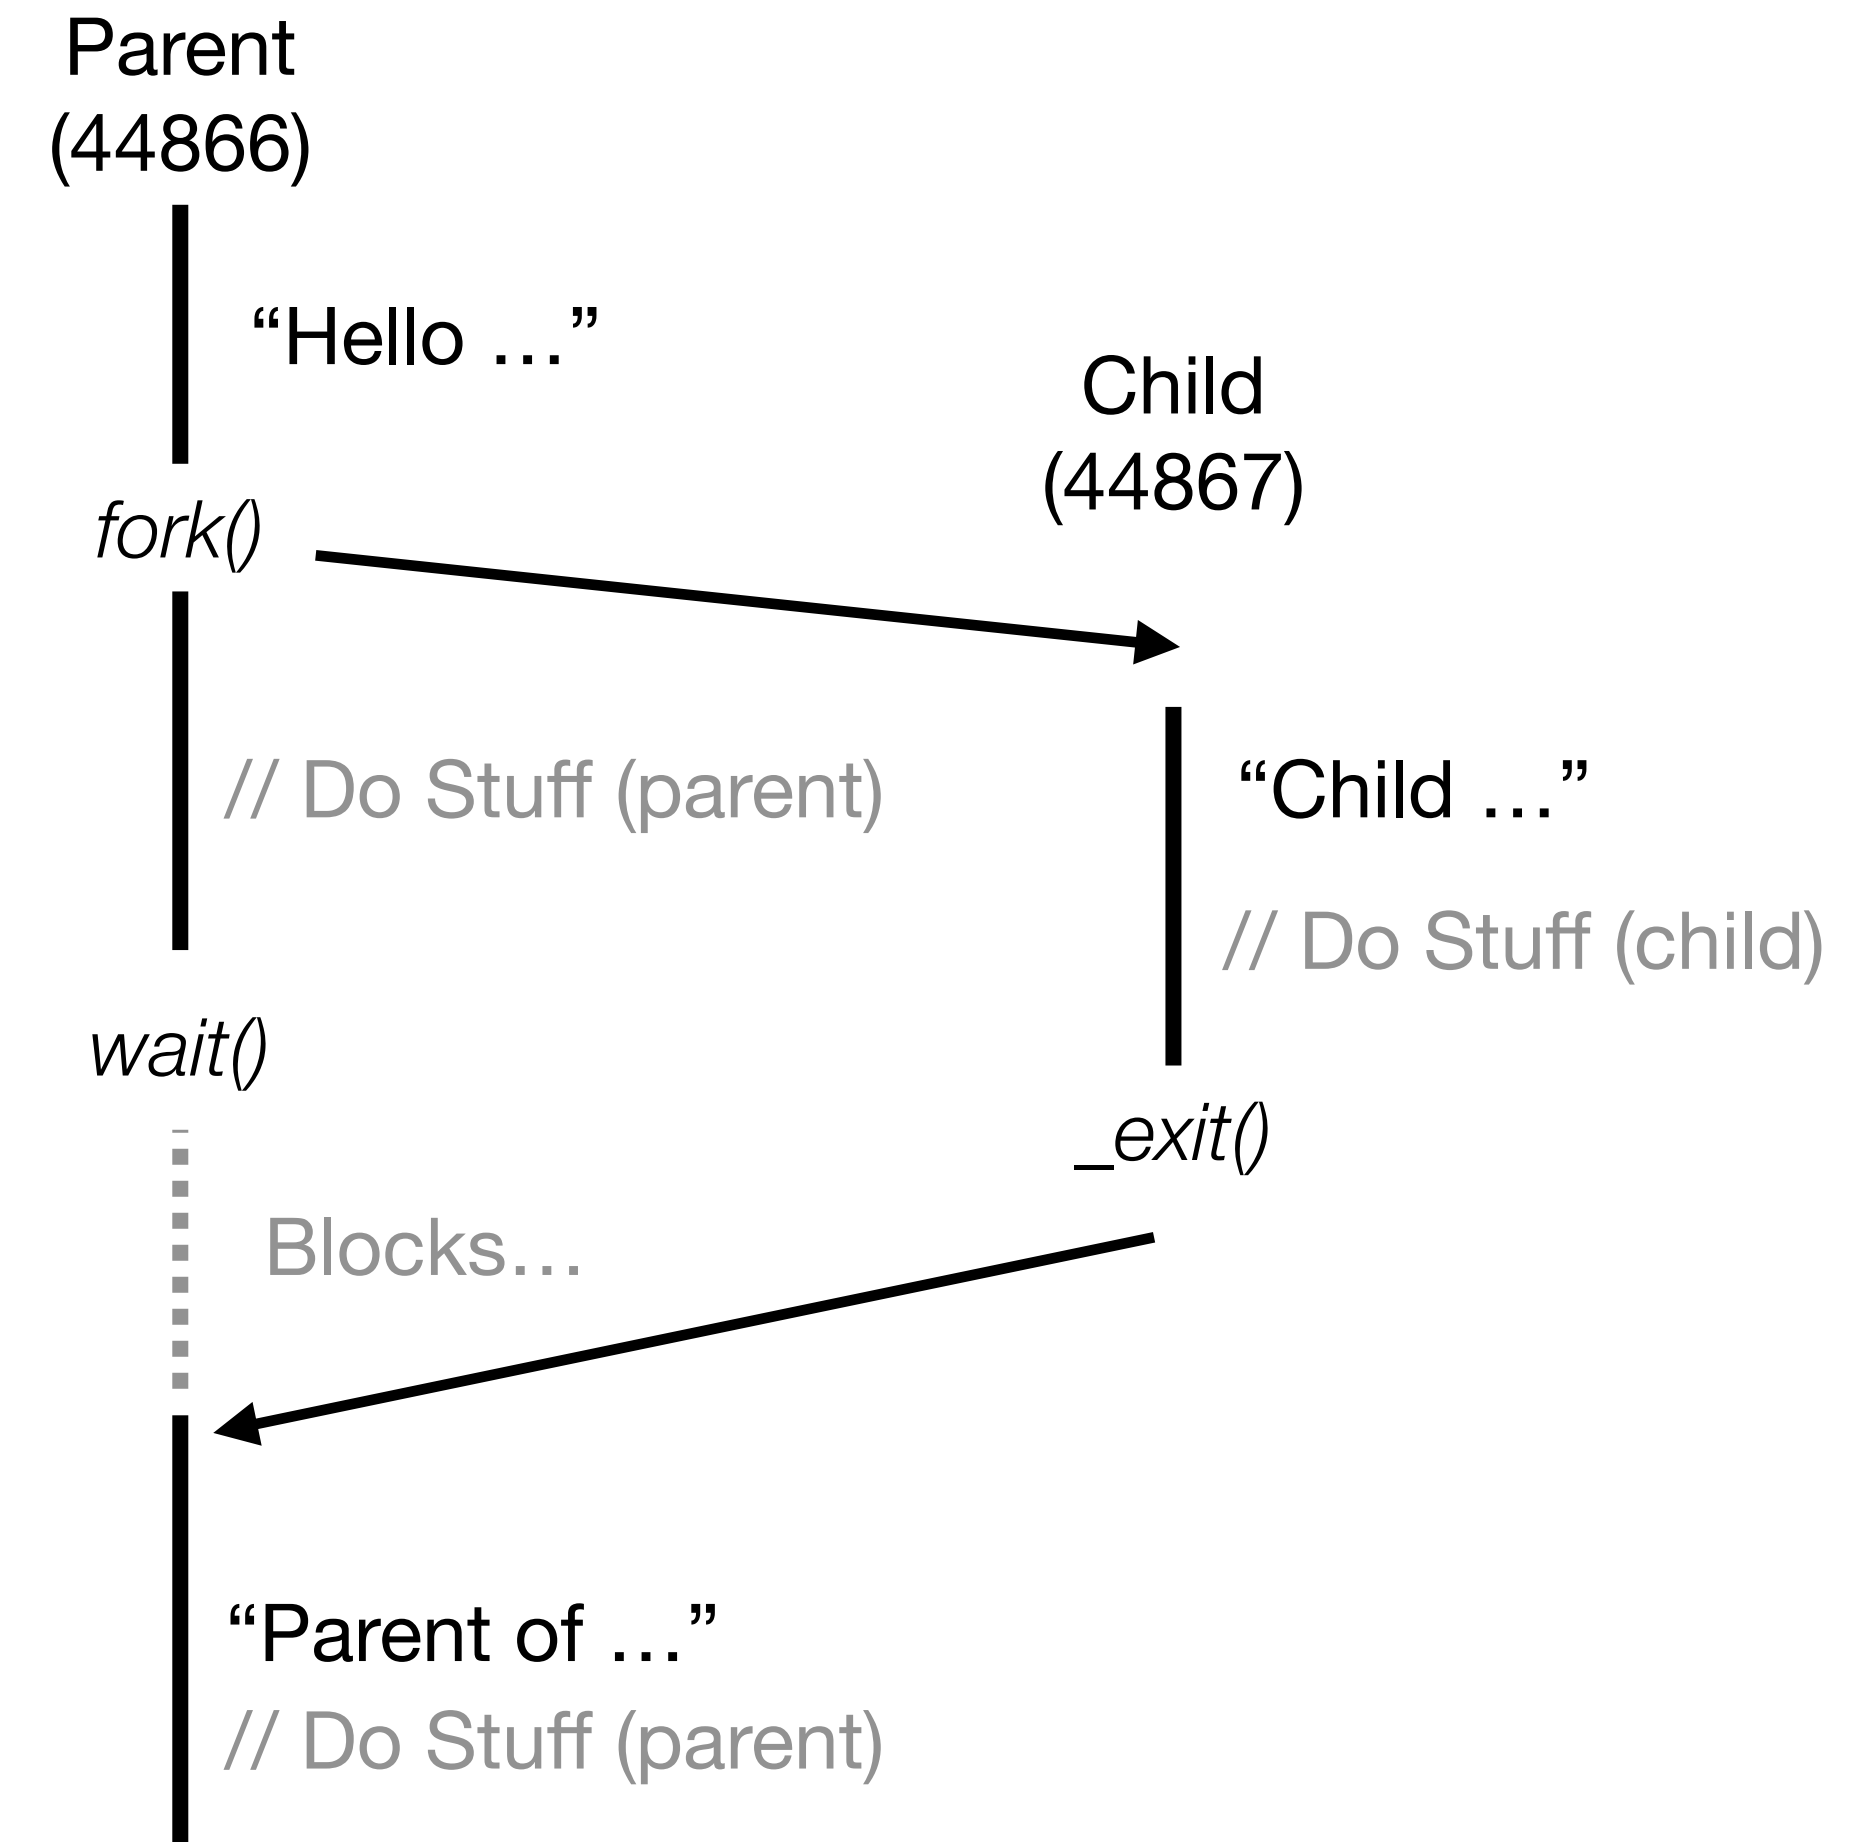
```
1 #include <unistd.h>
2 #include <stdio.h>
3 #include <sys/wait.h>
4
5 int main(int argc, char* argv[]){
6
7     printf("Hello (pid:%d)\n", getpid());
8     pid_t rc_fork = fork();
9     if(rc_fork == 0){
10         printf("Child (pid: %d)\n", getpid());
11         // Do stuff (child)
12         _exit(1);
13     }else{
14         //Do stuff (parent)
15         pid_t rc_wait = wait(NULL);
16         printf("Parent of %d (rc_wait: %d) (pid: %d)\n", rc_fork, rc_wait, getpid());
17         //Do stuff (parent)
18     }
19     return 0;
20 }
```

```
prompt> ./p1
Hello (pid:44866)
Child (pid: 44867)
Parent of 44867 (rc_wait: 44867) (pid: 44866)
prompt>
```

Wait example

```
1 #include <unistd.h>
2 #include <stdio.h>
3 #include <sys/wait.h>
4
5 int main(int argc, char* argv[]){
6
7     printf("Hello (pid:%d)\n", getpid());
8     pid_t rc_fork = fork();
9     if(rc_fork == 0){
10         printf("Child (pid: %d)\n", getpid());
11         // Do stuff (child)
12         _exit(1);
13     }else{
14         //Do stuff (parent)
15         pid_t rc_wait = wait(NULL);
16         printf("Parent of %d (rc_wait: %d) (pid: %d)\n", rc_fork, rc_wait, getpid());
17         //Do stuff (parent)
18     }
19     return 0;
20 }
```

```
prompt> ./p1
Hello (pid:44866)
Child (pid: 44867)
Parent of 44867 (rc_wait: 44867) (pid: 44866)
prompt>
```

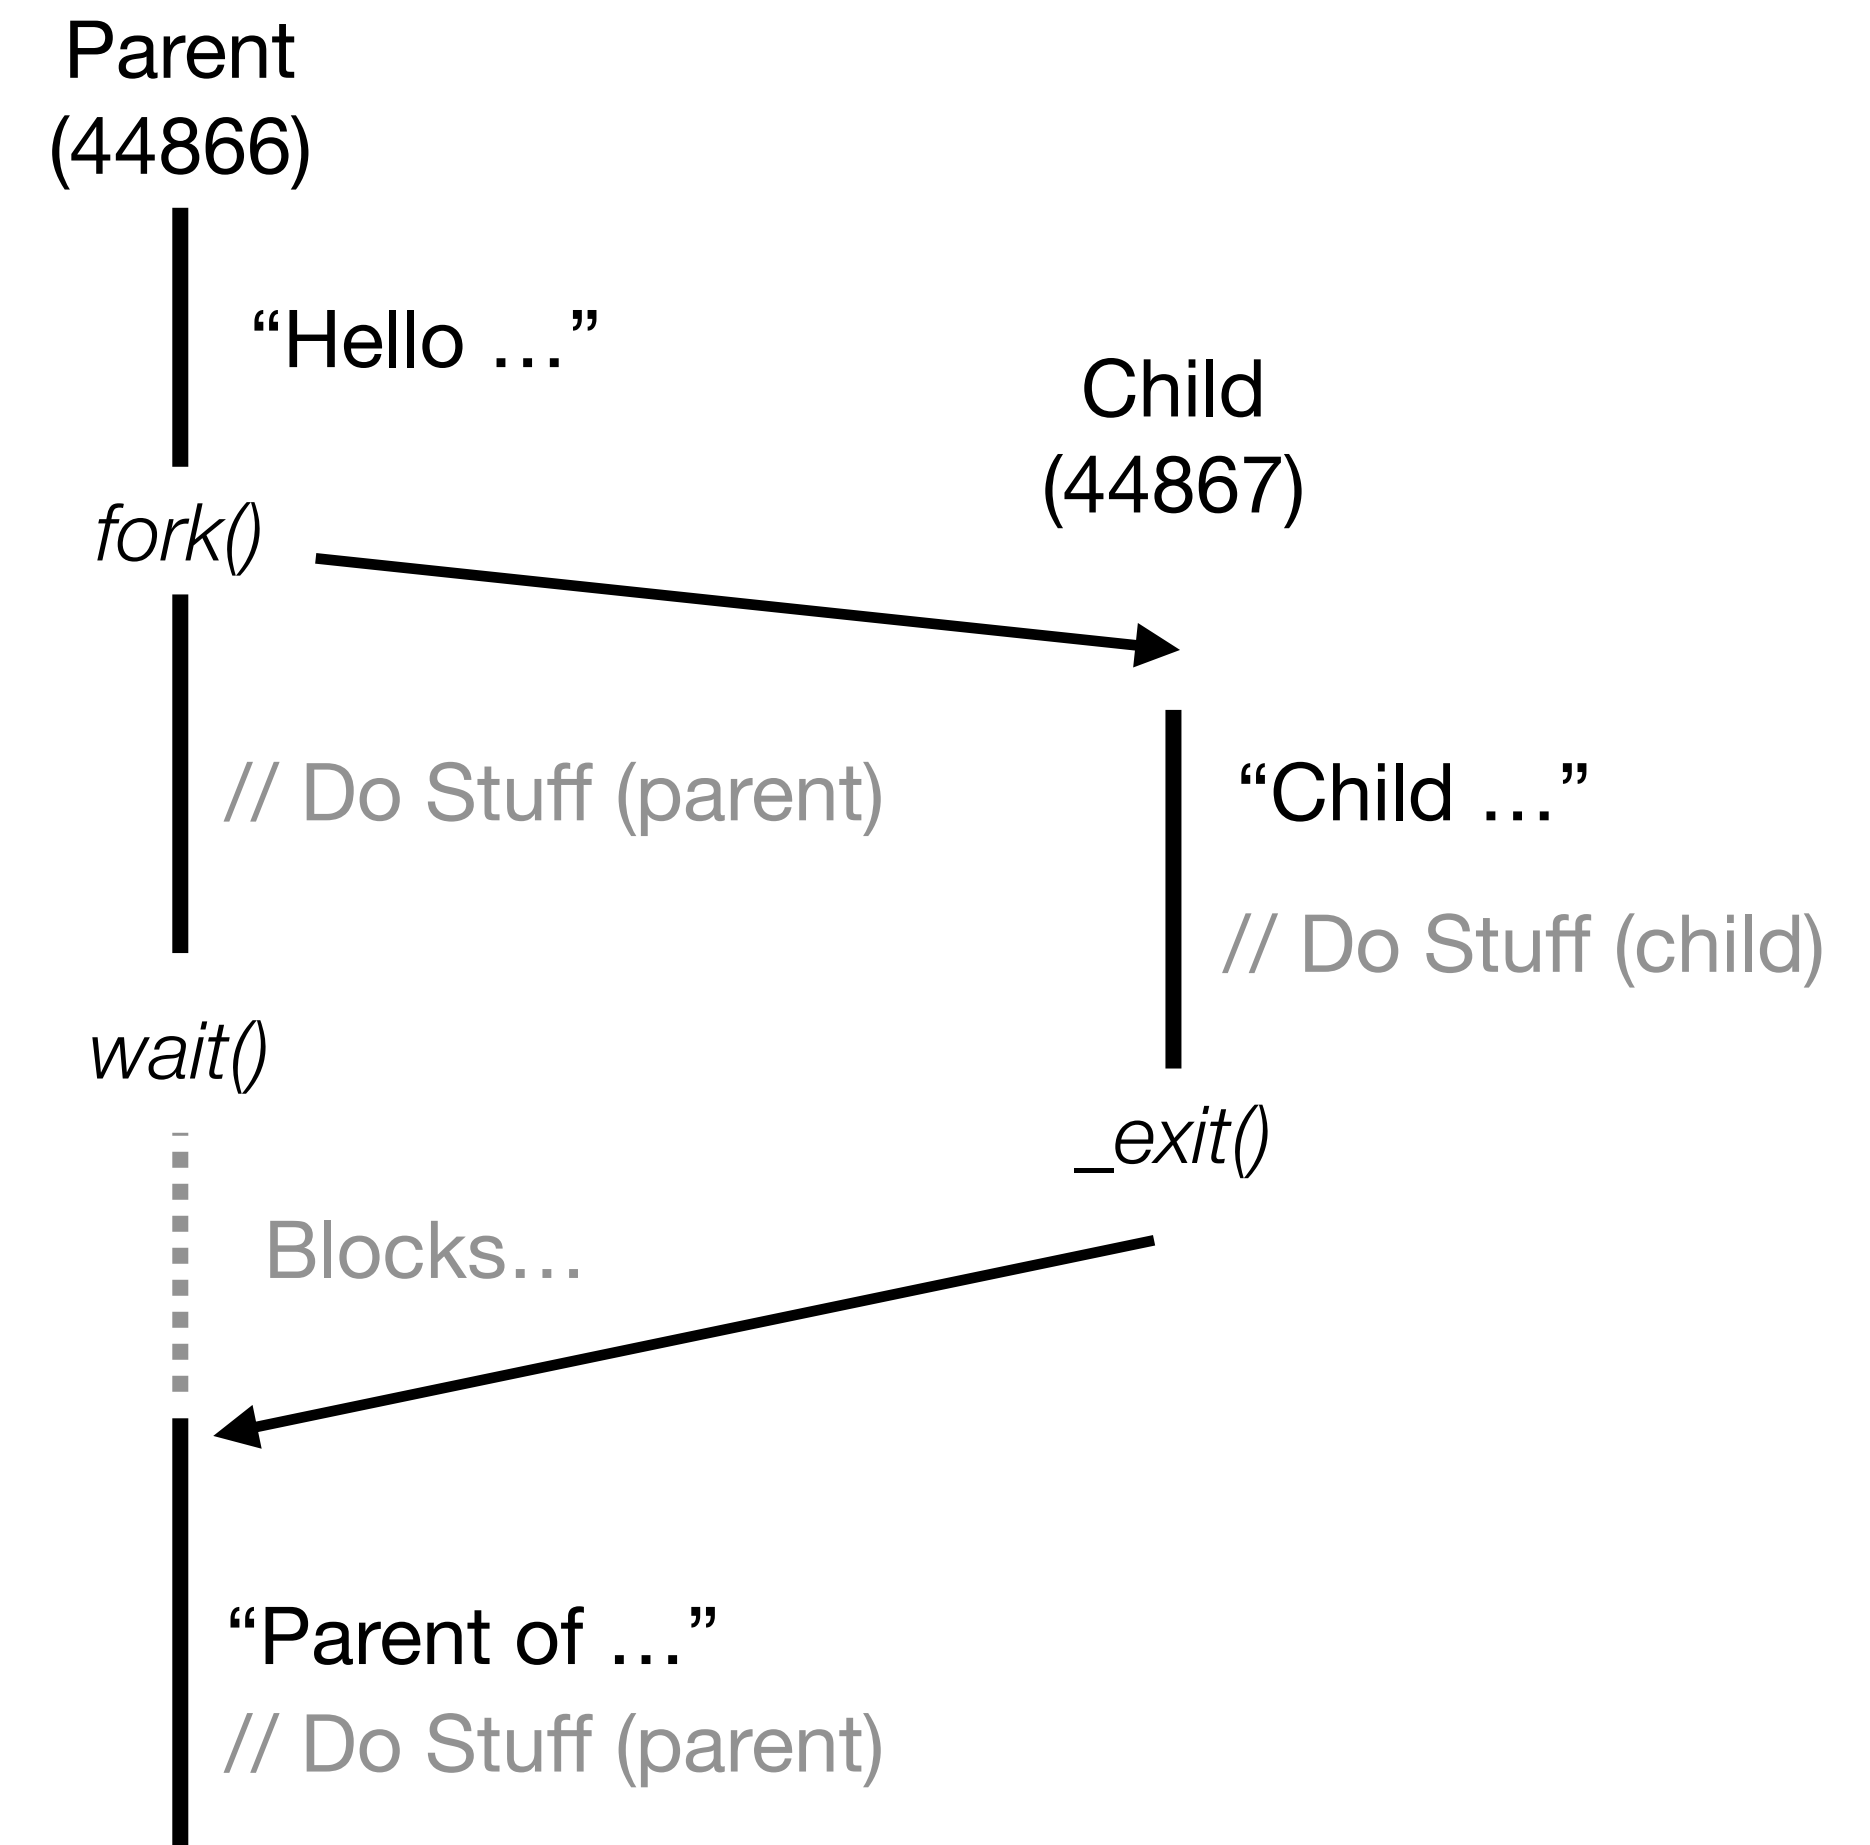


Fork and Wait

Concurrency

- © **Question:** If the previous program runs multiple times, the messages “*Child ...*” and “*Parent ...*” are printed always in the same order. Why?

```
prompt> ./p1
Hello (pid:44866)
Child (pid: 44867)
Parent of 44867 (rc_wait: 44867) (pid: 44866)
prompt>
```



Other details

Interesting things to see at the practical labs (Pt1)

- Do not forget to handle errors (when system calls return -1...)!
- Look at the *wait()* man page
 - Explore the *status* argument and the *WIFEXITED* and *WEXITSTATUS* macros
 - Explore the *waitpid()* call
- What happens if the child exits before the parent calls *wait()*?
 - The child goes into a **zombie** state (waiting for the parent to collect its status)
- What happens if the parent exits while the child is still executing?
 - The child becomes **orphaned**, being adopted by another process.
For example, the *init process* (*i.e.*, the initial process ran by UNIX systems)

Other details

Interesting things to see at the practical labs (Pt 2)

- How many processes does the code below create?
- Can you draw the tree of processes?

PID
1945



```
—1→ int main(int argc, char* argv){  
  2   fork()  
  3   fork()  
  4   fork()  
  5 }
```


Other details

Interesting things to see at the practical labs (Pt 2)

- How many processes does the code below create?
- Can you draw the tree of processes?



```
1 int main(int argc, char* argv){  
2     fork()  
3     fork()  
4     fork()  
5 }
```



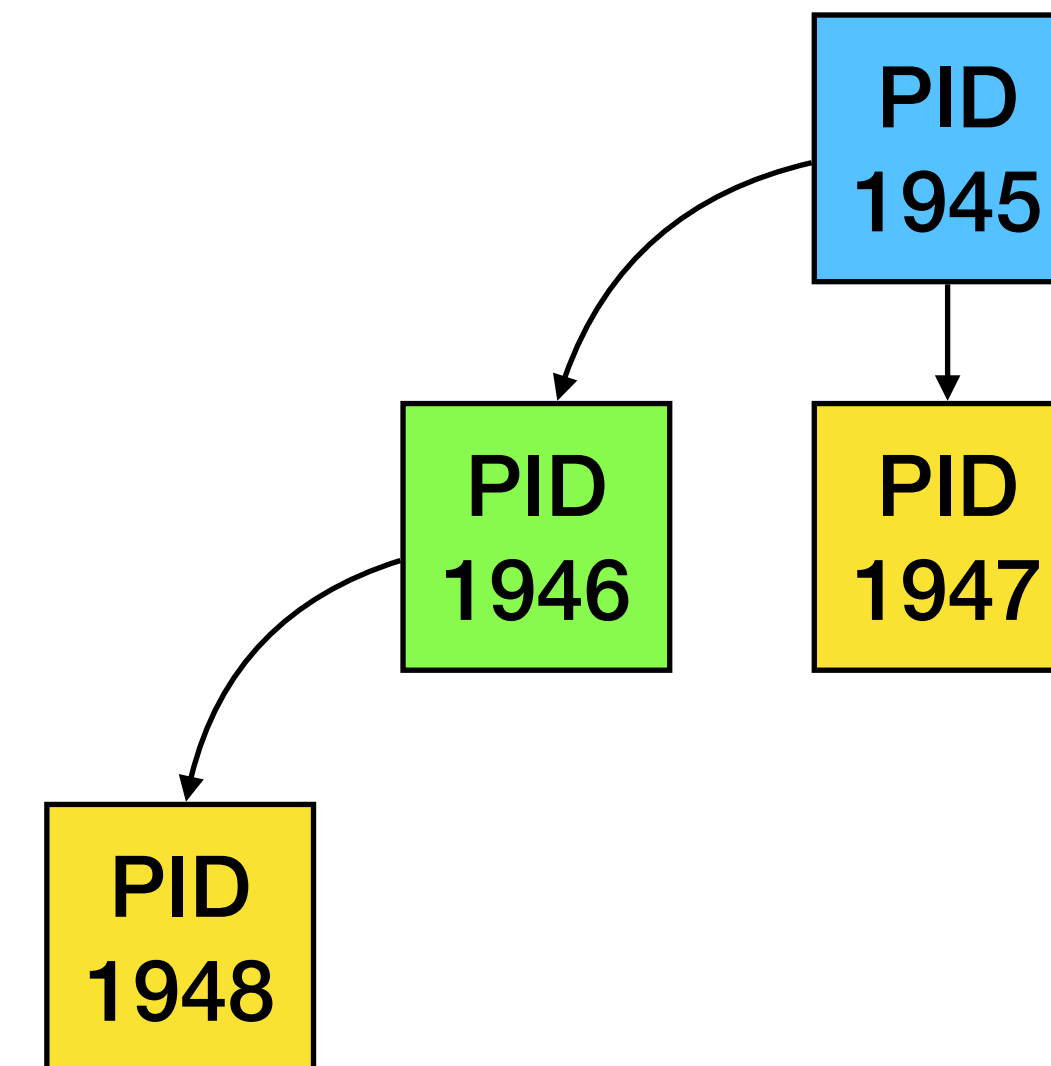
Other details

Interesting things to see at the practical labs (Pt 2)

- How many processes does the code below create?
- Can you draw the tree of processes?



```
1 int main(int argc, char* argv){  
2     fork()  
3     fork()  
4     fork()  
5 }
```




Other details

Interesting things to see at the practical labs (Pt 2)

- How many processes does the code below create?
- Can you draw the tree of processes?



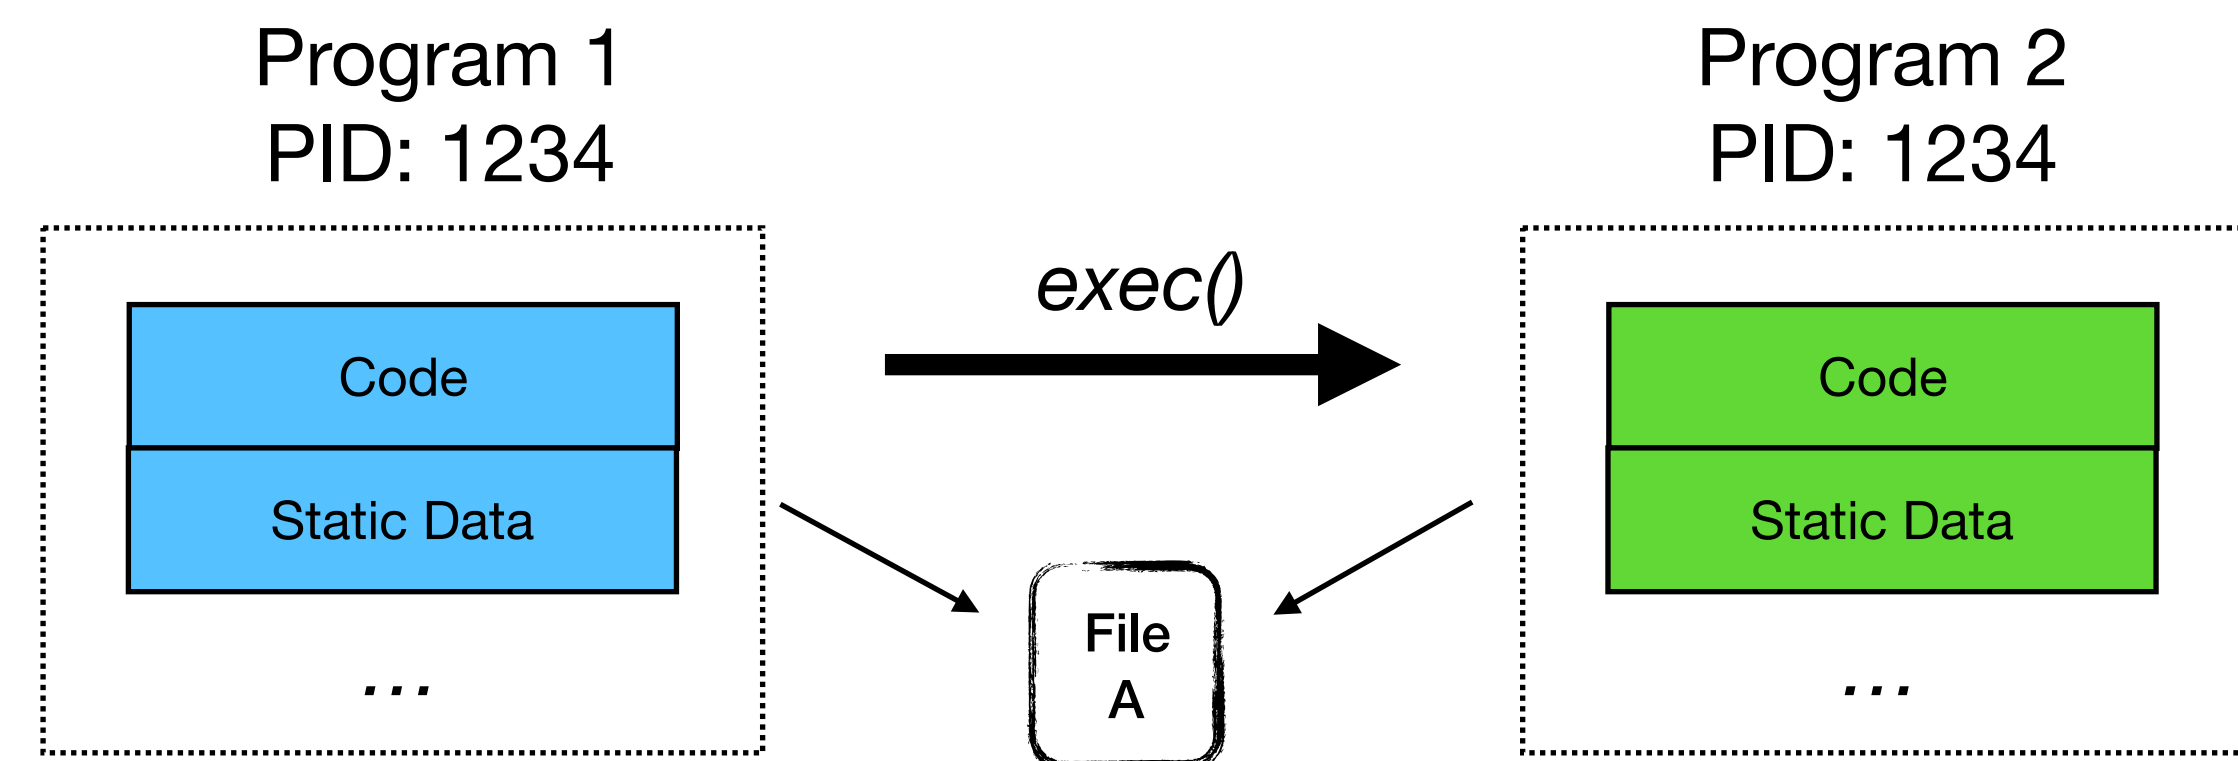
```
1 int main(int argc, char* argv){  
2     fork()  
3     fork()  
4     fork()  
5 }
```



UNIX Process API

Exec

- What if one wants the child process to run a different program?
- The `exec()` system call replaces the code and static data of the current process with the ones from a new program
 - Heap and stack are reinitialized
 - The new program starts at the `main()` entry point
 - PID is maintained (same process)
 - Initialized resources (e.g., opened files) are still usable by the process
- There are several variants of `exec` (e.g., `execv()`, `execp()`, `execvp()`, ...)



Exec example



```
1 #include <unistd.h>
2 #include <stdio.h>
3 #include <sys/wait.h>
4 #include <string.h>
5
6 int main(int argc, char* argv[]){
7
8     pid_t rc_fork = fork();
9     if(rc_fork == 0){
10         printf("Child (pid: %d)\n", getpid());
11         char *myargs[3];
12         myargs[0] = strdup("wc");          // program: wc
13         myargs[1] = strdup("exec.c");     // arg: file to count
14         myargs[2] = NULL;                 // mark end of array
15         → execvp(myargs[0], myargs);      // runs word count
16         printf("I'm Here Child (pid: %d)\n", getpid());
17         _exit(1);
18     }else{
19         pid_t rc_wait = wait(NULL);
20         printf("Parent of %d (pid: %d)\n", rc_fork, getpid());
21     }
22     return 0;
23 }
```

- `int execvp(const char *file, char *const argv[]);`
 - ▶ **file:** path/name of the program
 - ▶ **argv:** list of arguments of the program (terminated with NULL)
- The `wc` program executes on the child process
- **Question:** Is “I’m Here Child ...” printed in this program ?

Output of `wc` →

```
prompt> ./p1
Child (pid: 77262)
23   77  648 exec.c
Parent of 77262 (pid: 77261)
prompt>
```

Try running the following command in your shell: `wc exec.c`

Exec example



```
1 #include <unistd.h>
2 #include <stdio.h>
3 #include <sys/wait.h>
4 #include <string.h>
5
6 int main(int argc, char* argv[]){
7
8     pid_t rc_fork = fork();
9     if(rc_fork == 0){
10         printf("Child (pid: %d)\n", getpid());
11         char *myargs[3];
12         myargs[0] = strdup("wc");          // program: wc
13         myargs[1] = strdup("exec.c");      // arg: file to count
14         myargs[2] = NULL;                  // mark end of array
15         → execvp(myargs[0], myargs);       // runs word count
16         printf("I'm Here Child (pid: %d)\n", getpid());
17         _exit(1);
18     }else{
19         pid_t rc_wait = wait(NULL);
20         printf("Parent of %d (pid: %d)\n", rc_fork, getpid());
21     }
22     return 0;
23 }
```

◎ *int execvp(const char *file, char *const argv[]);*

- ▶ **file:** path/name of the program
- ▶ **argv:** list of arguments of the program (terminated with NULL)

◎ The wc program executes on the child process

◎ **Question:** Is “I’m Here Child ...” printed in this program ?

- ▶ exec() only returns to the original program if unsuccessful

Output of wc →

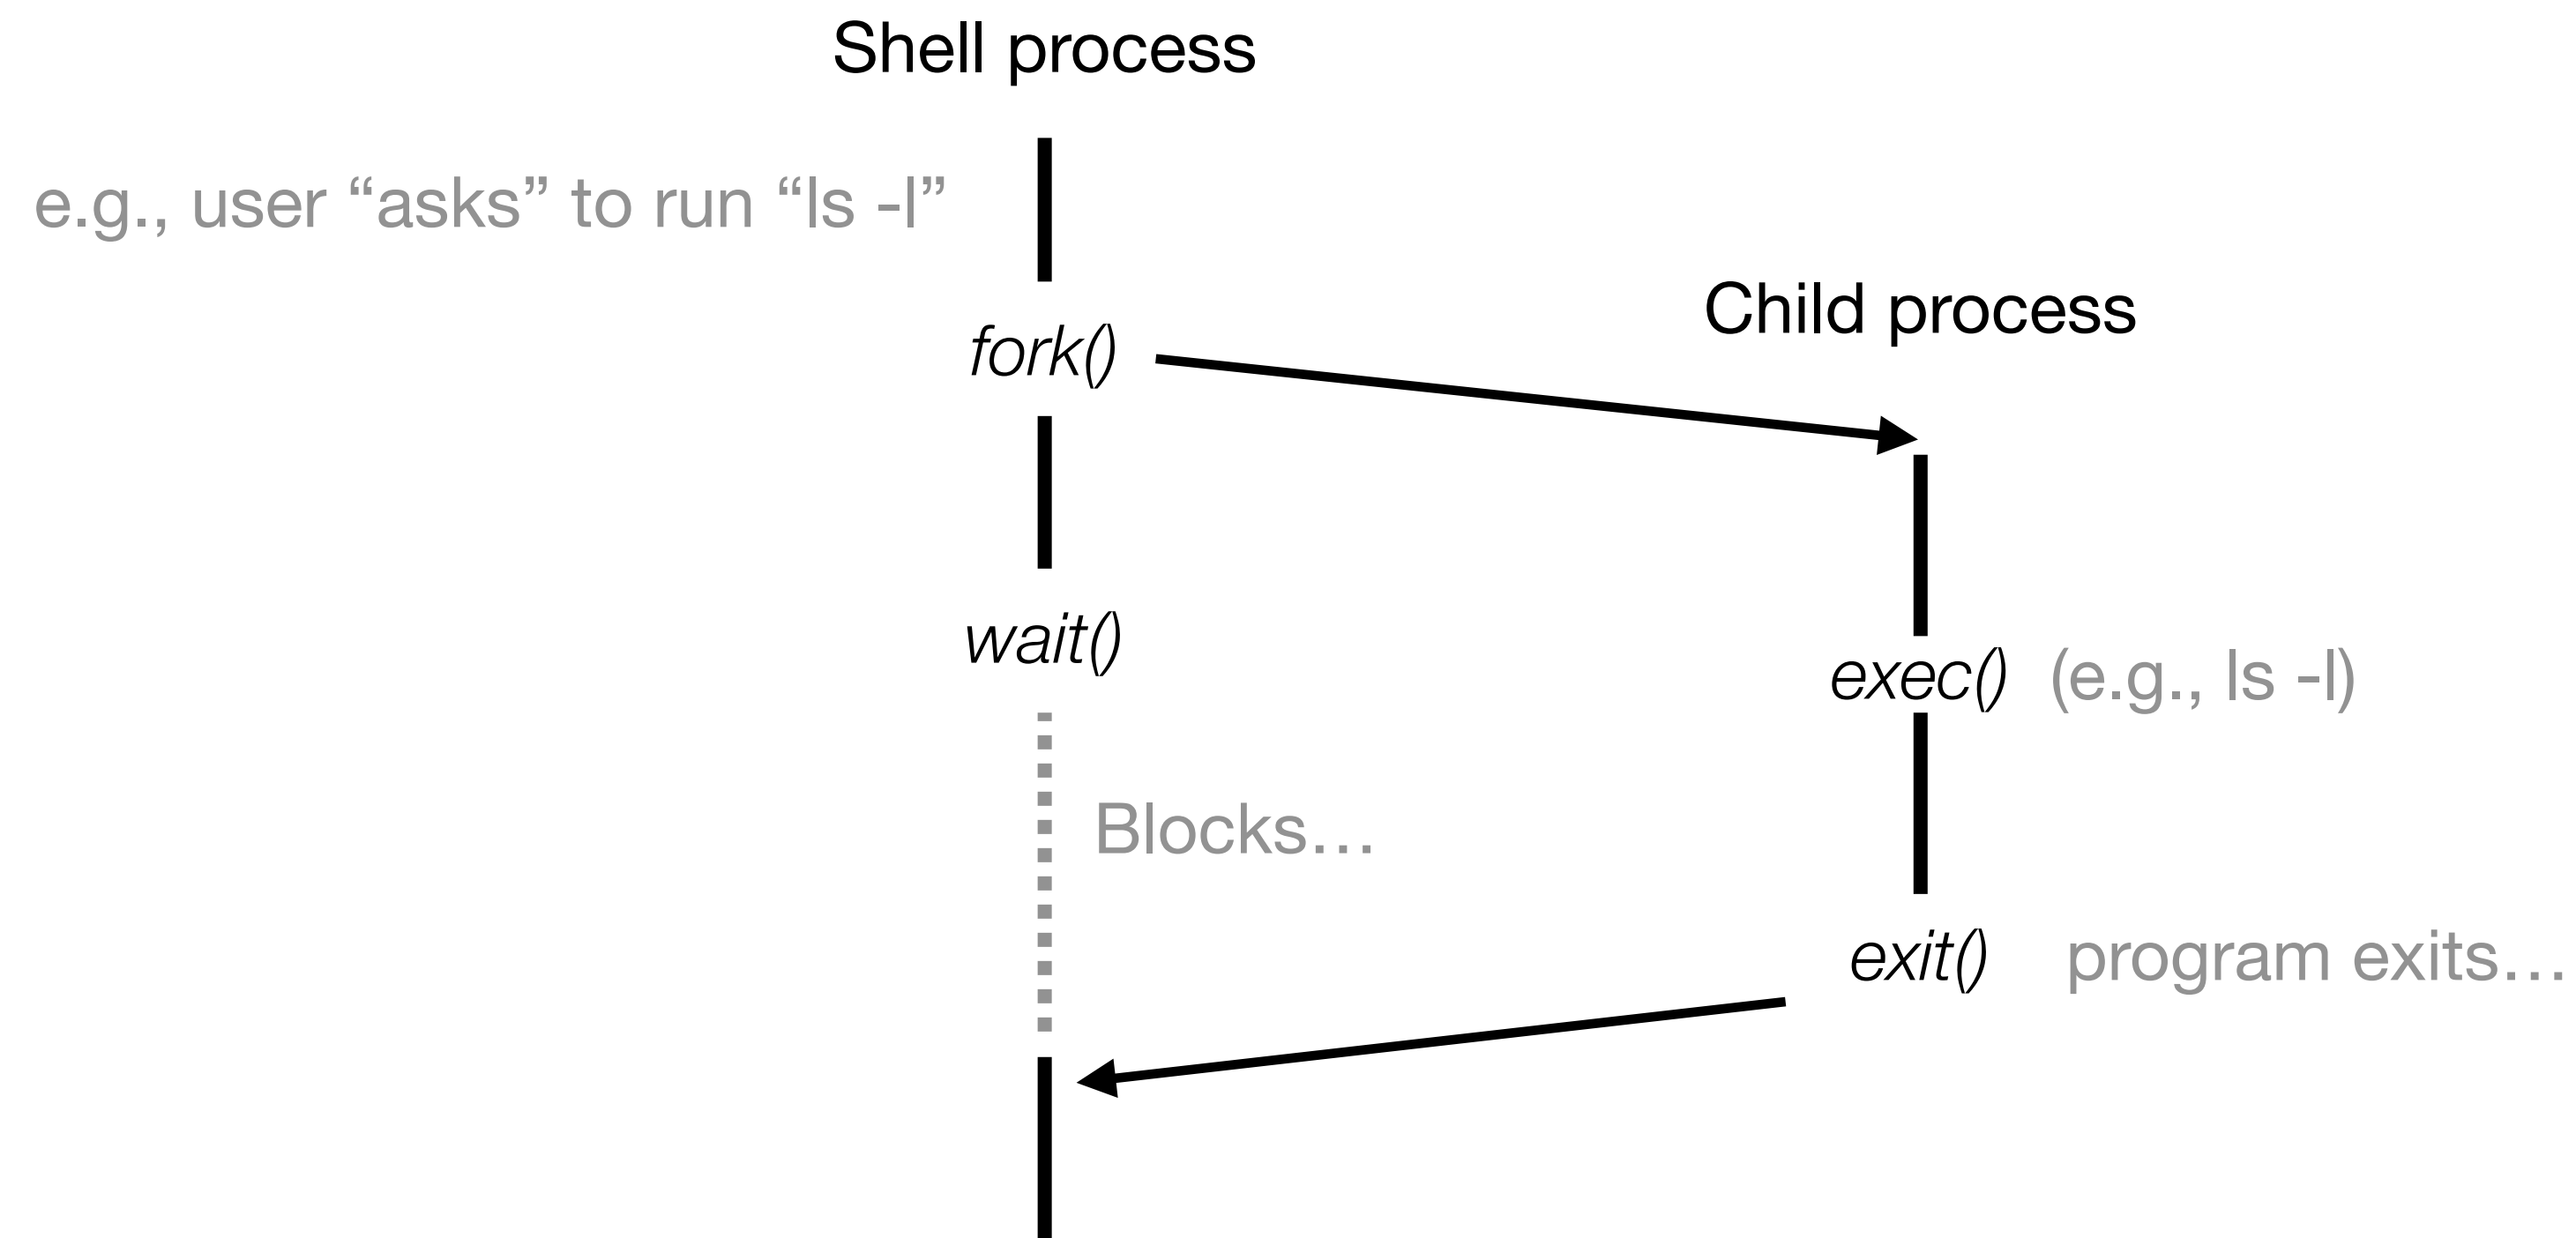
```
prompt> ./p1
Child (pid: 77262)
23   77  648 exec.c
Parent of 77262 (pid: 77261)
prompt>
```

Try running the following command in your shell: wc exec.c

Fork, Exec, Wait

A powerful combination


● **Question:** How does shell (bash) executes your programs?



Process Tree

In Unix and Linux

- In UNIX and Linux systems, processes are organized into a tree hierarchy
 - The first process to run after booting the OS is called **init**
 - All other processes descend from **init**
- Check the **ps** tool



Motivating the API

Why using `fork()` + `exec()` to run programs?

- Calling `fork()` and then `exec()` allows setting up the environment for new programs and to do interesting commands...
- **Example 1:** “`wc /etc/passwd > newfile.txt`”
 - Check the practical labs for the `dup()` system call
- **Example 2:** “`grep -v ^# /etc/passwd | wc`”
 - Check the practical labs for the `pipe()` and `dup()` system calls

Other system calls and tools

In Unix

- *kill()* system call used to
 - Pause, resume, force termination of processes
 - Configured as keystroke combinations in most UNIX shells
- Useful tools
 - **ps**, **pstree**, and **top** (check information about processes and system resources)
 - **kill** (same name as function but provided as a shell tool)

More Information

- **Chapters 4 and 5** - Remzi H. Arpaci-Dusseau, Andrea C. Arpaci-Dusseau. **Operating Systems: Three Easy Pieces**. Arpaci-Dusseau Books, 2018.
- Avi Silberschatz, Peter Baer Galvin, Greg Gagne. **Operating System Concepts (10. ed)**. John Wiley & Sons, 2018.

Questions?